

MATP 6600 Programming Project

Gabe Mancino-Ball

Fall 2018 - RPI

1 Problem

Given a set of training data $\{(\mathbf{x}_i, y_i)\}_{i=1}^m$ where $y_i \in \{-1, 1\}$ for each $i = 1, 2, \dots, m$, logistic regression can be used as a binary classifier. In this project, the methods of Newton, steepest gradient descent, and stochastic gradient descent are used on the datasets `gisette.mat` and `spamData.mat` to perform such classification.

2 Theoretical Approach

The model is formulated as

$$\min_{\mathbf{w}, b} f(\mathbf{w}, b) := \frac{1}{m} \sum_{i=1}^m \log(1 + \exp[-y_i(\mathbf{w}^T \mathbf{x}_i + b)]) + \frac{\lambda_1}{2} \|\mathbf{w}\|_2^2 + \frac{\lambda_2}{2} b^2 \quad (1)$$

where $\lambda_1 \geq 0, \lambda_2 \geq 0, b \in \mathbb{R}, y_i \in \{-1, 1\}$, and $\mathbf{x}_i, \mathbf{w} \in \mathbb{R}^n$ for $i = 1, 2, \dots, m$. Once a suitable $(\bar{\mathbf{w}}, \bar{b})$ have been found for (1), the hyperplane $\bar{\mathbf{w}}^T \mathbf{x}_{new} + \bar{b}$ is used to classify any new data.

2.1 Gradient and Hessian

The gradient and the Hessian matrix are necessary for the algorithms used. Using (1), the components of the gradient are formulated as

$$\begin{aligned} \frac{\partial}{\partial w_j} f &= \frac{1}{m} \sum_{i=1}^m \frac{-y_i x_{(i,j)}}{\exp[y_i(\mathbf{w}^T \mathbf{x}_i + b)] + 1} + \lambda_1 w_j \\ \frac{\partial}{\partial b} f &= \frac{1}{m} \sum_{i=1}^m \frac{-y_i}{\exp[y_i(\mathbf{w}^T \mathbf{x}_i + b)] + 1} + \lambda_2 b \end{aligned}$$

for $j = 1, 2, \dots, n$. Similarly, the components of the Hessian are formulated as

$$\begin{aligned} \frac{\partial^2}{\partial w_j \partial w_k} f &= \frac{1}{m} \sum_{i=1}^m \frac{x_{(i,j)} x_{(i,k)} \exp[y_i(\mathbf{w}^T \mathbf{x}_i + b)]}{(\exp[y_i(\mathbf{w}^T \mathbf{x}_i + b)] + 1)^2} \\ \frac{\partial^2}{\partial w_j \partial b} f &= \frac{\partial^2}{\partial b \partial w_j} = \frac{1}{m} \sum_{i=1}^m \frac{x_{(i,j)} \exp[y_i(\mathbf{w}^T \mathbf{x}_i + b)]}{(\exp[y_i(\mathbf{w}^T \mathbf{x}_i + b)] + 1)^2} \\ \frac{\partial^2}{\partial w_j^2} f &= \frac{1}{m} \sum_{i=1}^m \frac{x_{(i,j)}^2 \exp[y_i(\mathbf{w}^T \mathbf{x}_i + b)]}{(\exp[y_i(\mathbf{w}^T \mathbf{x}_i + b)] + 1)^2} + \lambda_1 \\ \frac{\partial^2}{\partial b^2} f &= \frac{1}{m} \sum_{i=1}^m \frac{\exp[y_i(\mathbf{w}^T \mathbf{x}_i + b)]}{(\exp[y_i(\mathbf{w}^T \mathbf{x}_i + b)] + 1)^2} + \lambda_2. \end{aligned}$$

3 Numerical Results

Using the method of steepest gradient descent, Newton’s method, and stochastic gradient descent, numerical experiments are implemented in MATLAB. See Listings (5), (6), and (7) for the code for these methods, respectively. Steepest gradient descent algorithm uses the method of *backtracking* coupled with Armijo’s Rule, which is outlined in [1], for determining the step size at each iteration; here $\sigma = 0.5$ and $\alpha_{new} = 0.75\alpha_{old}$. Pure Newton’s Method is used here (i.e. $\alpha = 1$). For the stochastic gradient descent method, a minibatch of 100 samples is taken for each iteration; the step size is given by $\alpha_k = \left(\frac{5}{6}\right)\frac{1}{k}$. Using the data sets `spamData.mat` and `gisette.mat` with the following termination rule:

$$\|\nabla(f(\mathbf{w}^{(k)}, b^{(k)}))\|_2 \leq \epsilon \max\{1, \|(\mathbf{w}^{(k)}, b^{(k)})\|_2\}$$

the algorithms are tested with $\epsilon \in \{10^{-2}, 10^{-4}, 10^{-6}\}$. For these experiments, $\lambda_1 = \lambda_2 = 0.001$ (see section 3.2 for results with varied λ_i , $i = 1, 2$). The initial guesses for $(\bar{\mathbf{w}}, \bar{b})$ are given by $(\mathbf{w}^{(0)}, b^{(0)}) = (\mathbf{0}, 0)$ for all of the experiments. The results are summarized in the following tables below.

Table 1: Results for the Spam dataset

| Method | Tolerance | Total Iteration Number | Total Run Time | Classification Accuracy |
|---------------|-----------|------------------------|----------------|-------------------------|
| Steepest GD | 10^{-2} | 10000 | 138.66 seconds | 87.19% |
| Newton’s | 10^{-2} | 7 | 0.5 seconds | 94.34% |
| Steepest GD | 10^{-4} | 25000 | 368.72 seconds | 90.18% |
| Newton’s | 10^{-4} | 9 | 0.64 seconds | 94.34% |
| Steepest GD | 10^{-6} | 50000 | 782.35 seconds | 91.18% |
| Newton’s | 10^{-6} | 9 | 0.67 seconds | 94.34% |
| Stochastic GD | - | 2000 | 8.5 seconds | 91.51% |

Table 2: Results for the Gisette dataset

| Method | Tolerance | Total Iteration Number | Total Run Time | Classification Accuracy |
|---------------|-----------|------------------------|----------------|-------------------------|
| Steepest GD | 10^{-2} | 183 | 8.21 seconds | 93.3% |
| Newton’s | 10^{-2} | 6 | 13.07 seconds | 93.4% |
| Steepest GD | 10^{-4} | 952 | 27.85 seconds | 93.4% |
| Newton’s | 10^{-4} | 10 | 25.05 seconds | 93.3% |
| Steepest GD | 10^{-6} | 5814 | 128.2 seconds | 93.3% |
| Newton’s | 10^{-6} | 11 | 26.25 seconds | 93.3% |
| Stochastic GD | - | 5000 | 136.87 seconds | 92% |

Note

The functions in Listings (1), (2), and (3) had originally been made with nested `for` loops; thanks to Robben Teufel, they are now vectorized. The MATLAB file `RunMe.m` will output the results for $\epsilon = 10^{-2}$ in about 5 minutes.

3.1 Plots of Objective Function

The following plots visually display how (1) is decreasing with each iteration with respect to the given method. Each plot shows the log of the difference between the computed objective value and the optimal objective value which is computed with Newton’s method (with $\epsilon = 10^{-6}$). The plots below only show the results for $\epsilon = 10^{-6}$, since these plots will be the most informative.

Figure 1: Comparison of Algorithms on Spam Data

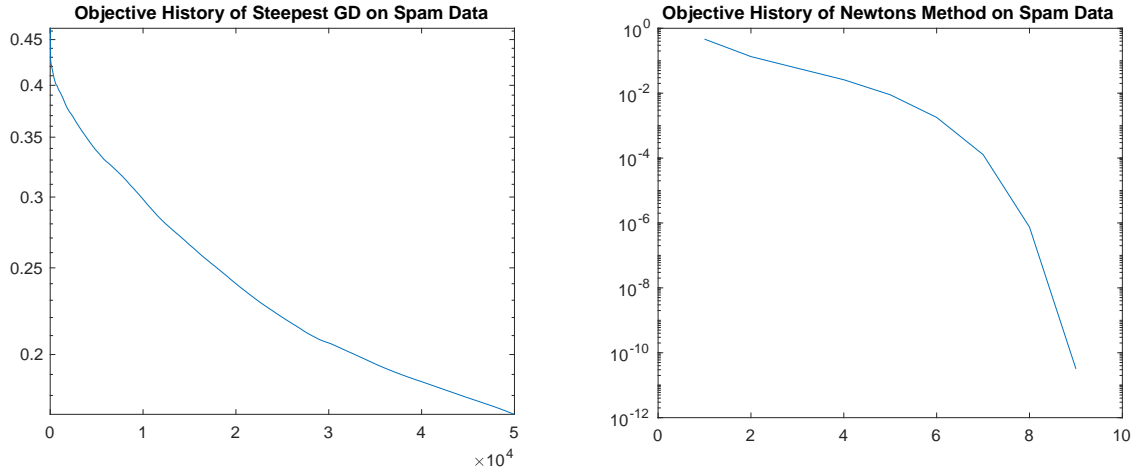


Figure 2: Comparison of Algorithms on Gisette Data

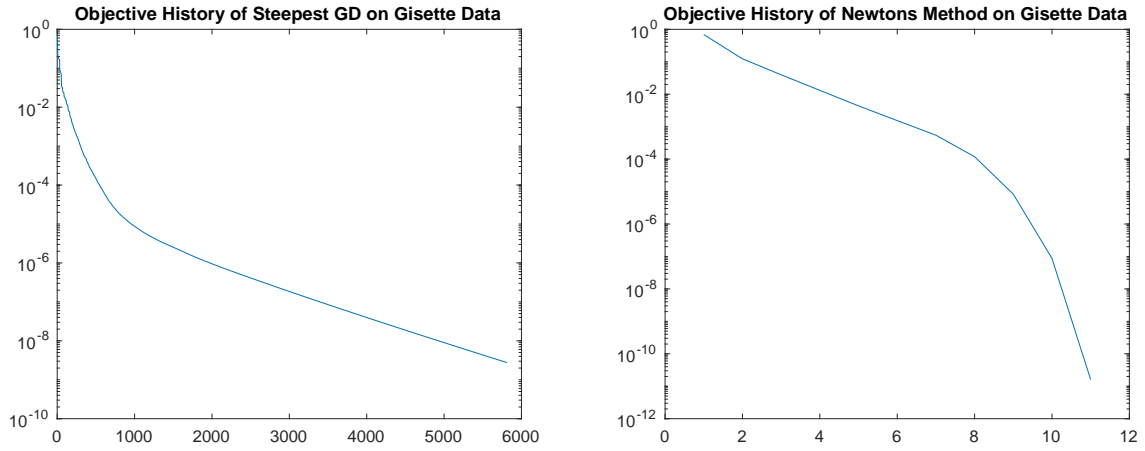
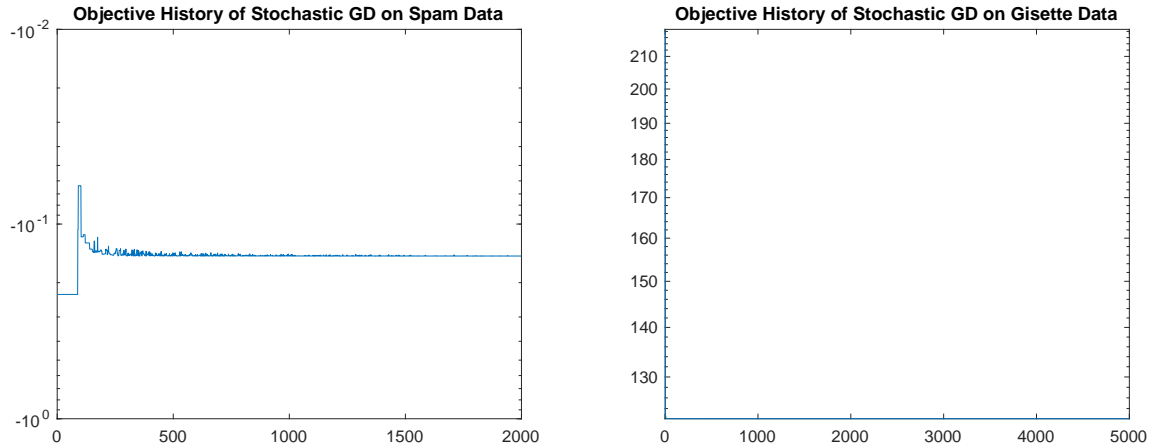


Figure 3: Stochastic Gradient Descent on each Dataset



3.2 Varying λ_1 and λ_2 : Impact on Classification Accuracy

The parameters λ_i , $i = 1, 2$, are varied in this section for the gradient descent method and Newton's method. For each test, the initial guess of $(\mathbf{w}^{(0)}, b^{(0)}) = (\mathbf{0}, 0)$ is made along with a tolerance of $\epsilon = 10^{-2}$. The resulting change in classification accuracy is summarized in the tables below.

Table 3: Results for the Spam dataset

| Method | λ_1 | λ_2 | Run Time | Classification Accuracy |
|-------------|-------------|-------------|---------------|-------------------------|
| Steepest GD | 0.1 | 0.1 | 45.48 seconds | 84.52% |
| Newton's | 0.1 | 0.1 | 0.5 seconds | 91.51% |
| Steepest GD | 0.005 | 0.05 | 40.6 seconds | 84.52% |
| Newton's | 0.005 | 0.05 | 0.43 seconds | 93.84% |
| Steepest GD | 1 | 0.0075 | 45.33 seconds | 84.02% |
| Newton's | 1 | 0.0075 | 0.53 seconds | 79.36% |

Table 4: Results for the Gisette dataset

| Method | λ_1 | λ_2 | Run Time | Classification Accuracy |
|-------------|-------------|-------------|---------------|-------------------------|
| Steepest GD | 0.1 | 0.1 | 9.28 seconds | 93.1% |
| Newton's | 0.1 | 0.1 | 14.53 seconds | 93.4% |
| Steepest GD | 0.005 | 0.05 | 8.5 seconds | 93.3% |
| Newton's | 0.005 | 0.05 | 12.96 seconds | 93.6% |
| Steepest GD | 1 | 0.0075 | 7.55 seconds | 92.3% |
| Newton's | 1 | 0.0075 | 11.63 seconds | 92.3% |

Using the tables above, it is clear that subtle changes in λ_1 and λ_2 had a greater impact on classification accuracy of the Spam dataset, for both steepest gradient descent and Newton's method. The Gisette dataset was more resilient to the affect of changing λ_1 and λ_2 . For plots of how (1) is changing, see Figure (3) at the end of this document.

4 Observations and Conclusion

In terms of performance, Newton's method was the superior algorithm; setting $\epsilon = 10^{-16}$ yielded a numerically optimal solution to (1) for both data sets. The steepest gradient descent algorithm did not converge for the Spam dataset, even after 50,000 iterations, this could be due to the step size at each iteration. For the stochastic gradient descent, the maximum number of iterations was set to a low value since the testing accuracy was above 90% for all trials. Overall, the methods have their advantages and disadvantages, but a large part of this is due to the given dataset. Comparison of algorithms, as done in this project, is necessary for determining an optimal algorithm to use on new data.

Note

Please note that the **Warning: Negative data ignored** comes from plotting the objective values on a logarithmic scale.

References

- [1] Mokhtar Bazaraa, Hanif Sherali, and C. M. Shetty. *Nonlinear Programming: Theory and Algorithms*. John Wiley & Sons, Inc., 2006.

MATLAB Code

Listing 1: Code for Objective Function

```
1 %-----
2 % Objective Function for LR
3 % (vectorized)
4 %
5 % Inputs:
6 % X(i,:) - ith data point
7 % y - vector of classification results
8 % w - normal vector to hyperplane
9 % b - scalar in hyperplane equation
10 % lambda1 - tuning parameter
11 % lambda2 - tuning parameter
12 %
13 % Outputs:
14 % f - function value
15 %-----
16 function f = UpdatedObjLR(X, y, w, b, lambda1, lambda2)
17     [m,n] = size(X);
18     f = 1/m*(sum(log(1+exp(-y.*(X*w+b))))) + 0.5*lambda1*(w'*w) + 0.5*lambda2*b^2;
19 end
```

Listing 2: Code for Gradient of Objective Function

```
1 %-----
2 % Gradient of Objective Function for LR
3 % (vectorized)
4 %
5 % Inputs:
6 % X(i,:) - ith data point
7 % y - vector of classification results
8 % w - normal vector to hyperplane
9 % b - scalar in hyperplane equation
10 % lambda1 - tuning parameter
11 % lambda2 - tuning parameter
12 %
13 % Outputs:
14 % gradf - gradient of function evaluated at w and b
15 %-----
16 function gradf = UpdatedGradLR(X, y, w, b, lambda1, lambda2)
17     [m,n] = size(X);
18     gradf = zeros(n+1,1);
19     gradf(1:n,1) = 1/m*sum((-y.*X./(exp(y.*(X*w+b))+1)))' + lambda1*w;
20     gradf(n+1,1) = 1/m*sum((-y./(exp(y.*(X*w+b))+1))) + lambda2*b;
21 end
```

Listing 3: Code for Hessian of Objective Function

```

1  %-----
2  % Hessian of Objective Function for LR
3  % (vectorized)
4  %
5  % Inputs:
6  % X(i,:) - ith data point
7  % y - vector of classification results
8  % w - normal vector to hyperplane
9  % b - scalar in hyperplane equation
10 % lambda1 - tuning parameter
11 % lambda2 - tuning parameter
12 %
13 % Outputs:
14 % H - hessian matrix of f
15 %-----
16 function H = UpdatedHessLR(X, y, w, b, lambda1, lambda2)
17     [m,n] = size(X);
18     H = zeros(n+1,n+1);
19     evec = exp(y.*(X*w+b))./(1+exp(y.*(X*w+b))).^2;
20     diage = diag(evec);
21     H(1:n,1:n) = (1/m)*(X'*diage)*X + eye(n)*lambda1;
22     H(1:n,n+1) = (1/m)*sum(X.*evec)';
23     H(n+1,1:n) = H(1:n,n+1)';
24     H(n+1,n+1) = (1/m)*sum(evec) + lambda2;
25 end

```

Listing 4: Code for Classifying New Data

```

1  %-----
2  % Classification Function for Logistic Regression
3  %
4  % Inputs:
5  % X(i,:) - ith data point
6  % w - normal vector for hyperplane
7  % b - scalar for hyperplane
8  %
9  % Outputs:
10 % y - vector with -1 and 1 as components
11 %-----
12 function y = ClassLR(X, w, b)
13     [m,n] = size(X);
14     y = zeros(m,1);
15     for i=1:m
16         if X(i,:)*w+b >= 0
17             y(i,1) = 1;
18         else
19             y(i,1) = -1;
20         end
21     end
22 end

```

Listing 5: Steepest Gradient Descent

```

1  %-----
2  % Steepest Gradient Descent for Logistic Regression
3  %
4  % Inputs:
5  % X(i,:) - ith data point as a row vector
6  % y - {-1, +1} classifier
7  % w - initial guess for w
8  % b - initial guess for b
9  % lambda1 - tuning parameter
10 % lambda2 - tuning parameter
11 % maxit - max number of iteration
12 % tol - tolerance
13 %
14 % Outputs:
15 % w - normal vector for hyperplane
16 % b - scalar for hyperplane
17 % hist_obj - history of objective value
18 % iter - number of iterations
19 %-----
20 function [w, b, iter, hist_obj] =...
21     SteepGD(X, y, w, b, lambda1, lambda2, maxit, tol)
22     [m,n] = size(X);
23     iter = 1;
24     grad = UpdatedGradLR(X, y, w, b, lambda1, lambda2);
25     obj = UpdatedObjLR(X, y, w, b, lambda1, lambda2);
26     hist_obj = obj;
27     while iter < maxit && norm(grad(1:n,1))+norm(grad(n+1,1))...
28         >= tol*max(1, norm(w)+norm(b))
29
30         % Loop for alpha using backtracking
31         alpha = 1;
32         while UpdatedObjLR(X, y, w - alpha*grad(1:n,1),...
33             b - alpha*grad(n+1,1), lambda1, lambda2) - ...
34             UpdatedObjLR(X, y, w, b, lambda1, lambda2) >= 0.5*alpha*...
35             grad'*(-grad)
36             alpha = 0.75*alpha;
37         end
38
39         % Update w and b
40         w = w - alpha*grad(1:n,1);
41         b = b - alpha*grad(n+1,1);
42
43         % Update gradient
44         grad = UpdatedGradLR(X, y, w, b, lambda1, lambda2);
45
46         % Update objective
47         obj = UpdatedObjLR(X, y, w, b, lambda1, lambda2);
48         hist_obj = [hist_obj; obj];
49         iter = iter + 1;
50     end
51 end

```

Listing 6: Newton's Method

```

1  %-----
2  % Newton's Method for Logistic Regression
3  %
4  % Inputs:
5  % X(i,:) - ith data point as a row vector
6  % y - {-1, +1} classifier
7  % w - initial guess for w
8  % b - initial guess for b
9  % lambda1 - tuning parameter
10 % lambda2 - tuning parameter
11 % maxit - max number of iteration
12 % tol - tolerance
13 %
14 % Outputs:
15 % w - normal vector for hyperplane
16 % b - scalar for hyperplane
17 % hist_obj - history of objective value
18 % iter - number of iterations
19 %-----
20 function [w, b, iter, hist_obj] =...
21     Newton(X, y, w, b, lambda1, lambda2, maxit, tol)
22     [m,n] = size(X);
23     iter = 1;
24     grad = UpdatedGradLR(X, y, w, b, lambda1, lambda2);
25     H = UpdatedHessLR(X, y, w, b, lambda1, lambda2);
26     obj = UpdatedObjLR(X, y, w, b, lambda1, lambda2);
27     hist_obj = obj;
28     % Pure Newton's
29     alpha = 1;
30
31     while iter < maxit && norm(grad(1:n,1))+norm(grad(n+1,1))...
32         >= tol*max(1, norm(w)+norm(b))
33
34         % Find descent direction
35         d = linsolve(H, grad);
36
37         % Update w and b
38         w = w - alpha*d(1:n,1);
39         b = b - alpha*d(n+1,1);
40
41         % Update hessian and gradient
42         H = UpdatedHessLR(X, y, w, b, lambda1, lambda2);
43         grad = UpdatedGradLR(X, y, w, b, lambda1, lambda2);
44
45         % Update objective
46         obj = UpdatedObjLR(X, y, w, b, lambda1, lambda2);
47         hist_obj = [hist_obj; obj];
48         iter = iter + 1;
49     end
50 end

```


Listing 7: Stochastic Gradient Descent

```

1  %-----
2  % Stochastic Gradient Descent for Logistic Regression
3  %
4  % Inputs:
5  % X(i,:) - ith data point as a row vector
6  % y - {-1, +1} classifier
7  % w - initial guess for w
8  % b - initial guess for b
9  % lambda1 - tuning parameter
10 % lambda2 - tuning parameter
11 % maxit - max number of iteration
12 % tol - tolerance
13 % batch - batchsize
14 %
15 % Outputs:
16 % w - normal vector for hyperplane
17 % b - scalar for hyperplane
18 % hist_obj - history of objective value
19 % iter - number of iterations
20 %-----
21 function [w, b, iter, hist_obj] =...
22     SGD(X, y, w, b, lambda1, lambda2, maxit, batch)
23     [m,n] = size(X);
24     hist_obj = 0;
25     iter = 1;
26     while iter < maxit
27         for i=1:(m/batch)
28             % Pick random entries
29             r = randi([1,m], batch, 1);
30             Xup = X(r,:);
31             yup = y(r,1);
32             sgrad = UpdatedGradLR(Xup, yup, w, b, lambda1, lambda2);
33
34             % Set alpha
35             alpha = (5/6)/(iter);
36
37             % Update w and b
38             w = w - alpha*sgrad(1:n,1);
39             b = b - alpha*sgrad(n+1,1);
40
41         end
42         % Update objective only if obj != Inf
43         obj = UpdatedObjLR(X, y, w, b, lambda1, lambda2);
44         if obj~=Inf
45             hist_obj = [hist_obj; (sum(hist_obj) + obj)/iter];
46         else
47             hist_obj = [hist_obj; hist_obj(length(hist_obj))];
48         end
49         iter = iter + 1;
50     end
51 end

```

Figure 4: Comparison of Algorithms with Varied λ_i

